

Tokenization + Algebra: How Language Models Learn to Count, Classify, and Compute

By Francisco Revelles

Author's Note

This article represents a refined version reviewed by both **GPT-5** and **Gemini**, two AI evaluators used to test conceptual clarity and mathematical framing. It distinguishes analogy from mechanism, defines terminology precisely, and incorporates supporting references to contemporary theoretical work.

A Note on This Framework

The ideas presented here offer a **conceptual lens** for interpreting large language models through algebra.

This is an *analogy* intended to build intuition—not a literal description of algorithms or a proven mechanistic model.

Introduction

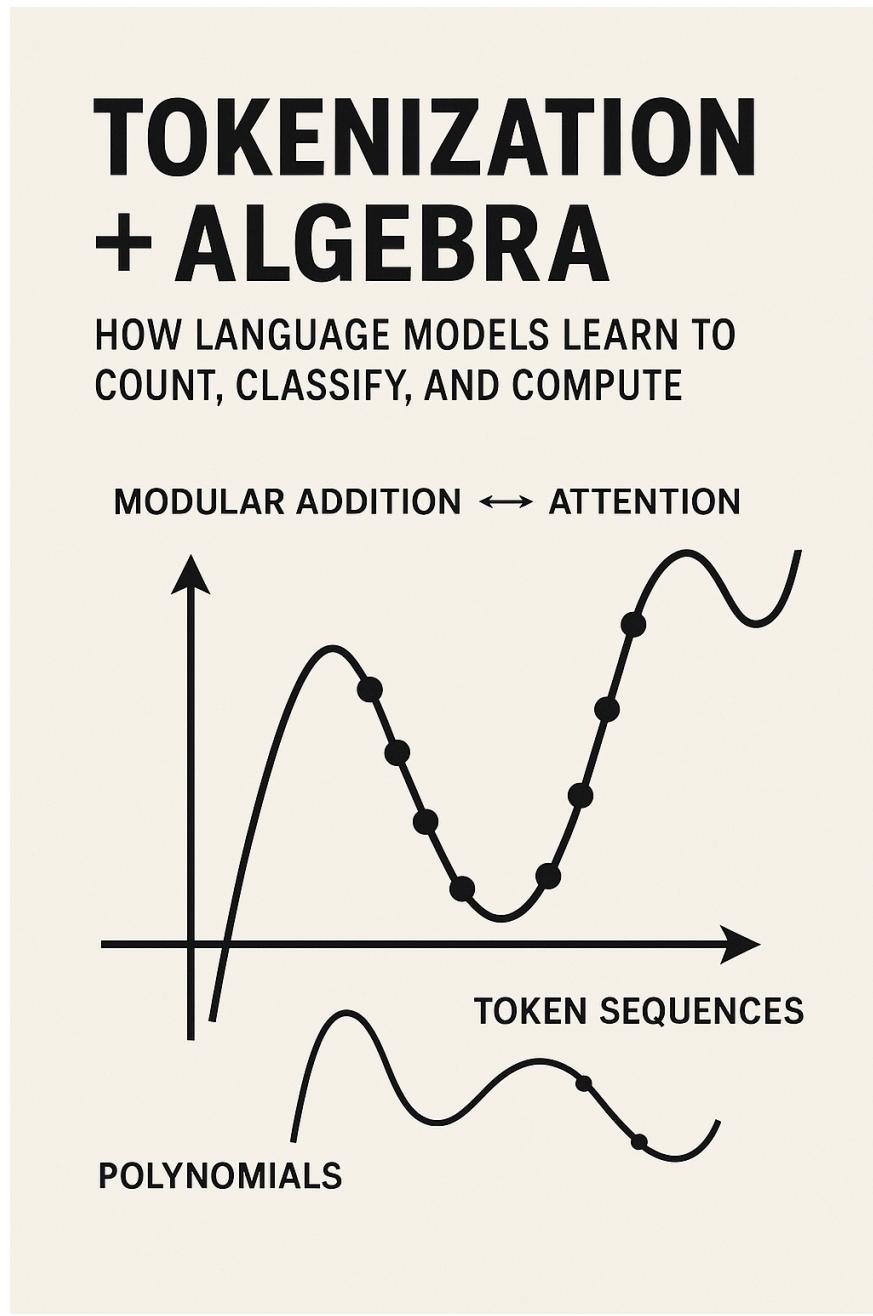


Figure 0—Cover illustration for Tokenization + Algebra

The graphic highlights the central parallels between algebra and language models: polynomials with token sequences, and modular addition with attention.

Language Models (LLMs) aren't just statistical parrots—they can also be understood as **algebraic machines**. Each token they process behaves like an element within a structured system of relationships, and the operations that connect those tokens resemble the

arithmetic logic behind **cryptography, signal processing, and number theory**.

Throughout this article, algebra serves as an interpretive lens connecting symbolic language units with the mathematical reasoning patterns that emerge in neural architectures.

In short, **tokenization is more than preprocessing—it is a gateway to algebra**. From token IDs to embeddings, from attention matrices to emergent reasoning, the architecture of LLMs can be viewed through the lens of **modular arithmetic**.

• • •

1. Tokenization as a Gateway to Algebra

Tokenization maps language into a finite alphabet of symbols—much like **elements of a finite field**, such as integers modulo a prime p (\mathbb{Z}_p).

- Tokens are **atomic**, each with a unique integer ID.
- Token embeddings are **vectors in continuous space**, yet their transformations behave in modular or categorical ways.
- Sequences of tokens resemble **polynomials over finite fields**, providing a natural analogy for reasoning about structure and composition.

For example, the mapping from tokens to IDs (see Figure 1) mirrors how elements in \mathbb{Z}_p index positions in modular arithmetic.

Seen this way, the building blocks of language are not arbitrary—they are **algebraic structures in disguise**.
2. Modular Arithmetic Inside LLMs

• • •

2. Modular Arithmetic Inside LLMs

Empirically, we cannot yet prove these parallels; however, as a **conceptual model**, modular arithmetic illuminates how information might circulate through a network.

As discussed in *How Modular Arithmetic Powers Everything from AI to Cryptography*, LLMs may perform operations that resemble algebraic computation:

- Hidden layers perform transformations that **may** correspond to modular addition or multiplication.
- Embedding spaces and attention matrices **may exhibit** periodic or cyclic behavior analogous to finite groups.
- The “wrap-around” behavior in transformers resembles arithmetic rollover, where values repeat modulo n .

These are not proven mechanisms—but they provide a powerful way to think about how symbolic structures persist through continuous spaces.³ The Algebraic View of Model Layers

• • •

3. The Algebraic View of Model Layers

We can visualize the model’s structure as a set of **algebraic lifts**—each stage translating discrete symbols into richer mathematical representations.

- **Helix metaphor:** imagine each token’s representation spiraling upward through layers like a helix, marking conceptual phase shifts in meaning.
- **Lexical layer:** discrete tokens and subwords (\mathbb{Z}_p).
- **Embedding layer:** continuous vector basis (\mathbb{R}^n).
- **Attention layers:** rotational transformations analogous to modular arithmetic.

- **Output projection:** grounding back into discrete language.

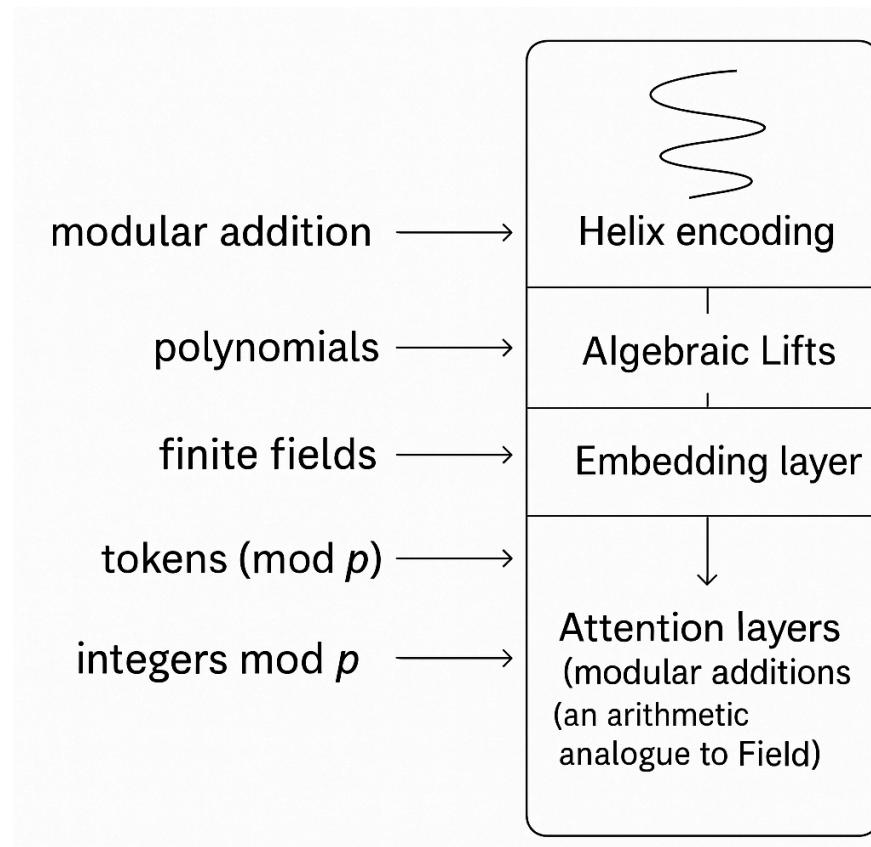


Figure 1—The Tokenization → Algebra Pipeline
Language tokens are lifted from \mathbb{Z}_p (discrete) into \mathbb{R}^n (continuous), transformed by modular-like rotations, and projected back into symbols.

... . . .

4. Practical Implications for LLM Design

This algebraic perspective has design consequences:

- **Tokenizer optimization:** Subword tokenizers can act like “prime factor bases,” reducing carry-over in sequence reasoning.
- **Algebra-aware architectures:** Embedding spaces constrained to algebraically inspired lattices could improve modular reasoning.
- **Error diagnosis:** Misalignments in arithmetic analogy often correlate with hallucination, suggesting new tools for

interpretability.

By treating language as algebra, we can design models that reason more robustly and fail more predictably.

TOKENIZATION + ALGEBRA

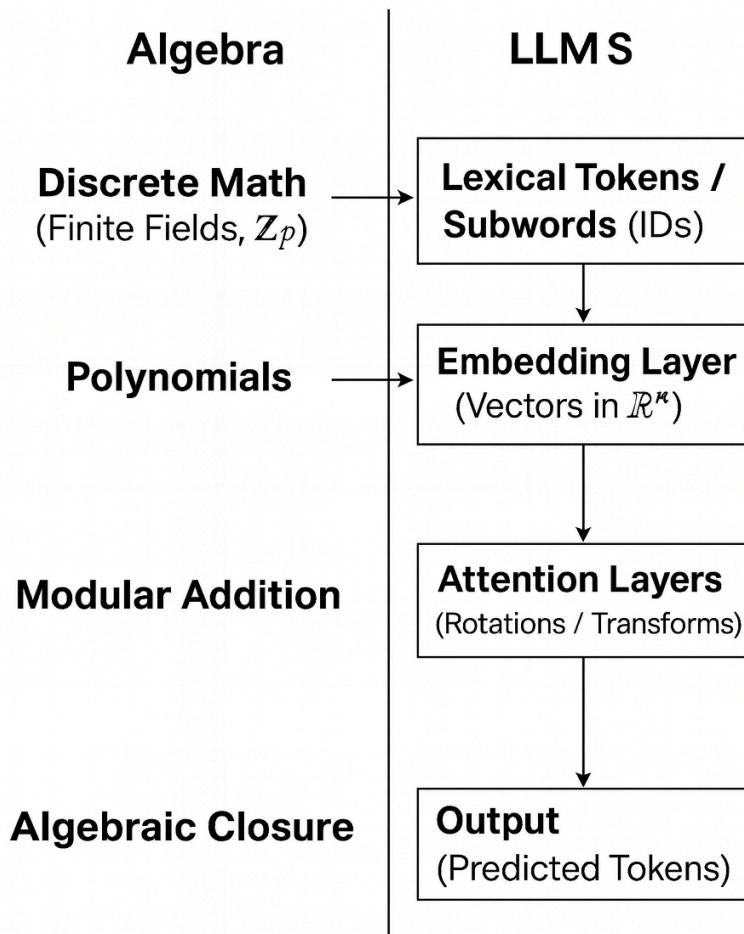


Figure 2. Attention as Modular Arithmetic.

This diagram illustrates how the internal mechanics of large language models (LLMs) mirror algebraic structures. Each stage in an LLM—from tokenization to attention—corresponds to a classical algebraic operation: discrete math (finite fields, Z_p) maps to tokenization, polynomials to embeddings (R^n), modular addition to attention transformations, and algebraic closure to output generation.

5. Toward a General Theory of Token Algebra

Across both theory and practice, these ideas point toward a unifying principle:

LLMs may be fruitfully interpreted as algebraic processors—discrete tokens lifted into continuous embeddings, transformed through modular rotations, and projected back into language.

This framing bridges:

- **Symbolic AI**—discrete logic and reasoning.
- **Deep Learning**—continuous optimization.
- **Pure Mathematics**—field theory and modular arithmetic.

By making this algebraic nature explicit, we open a path toward architectures that reason as elegantly as they generate.

The Algebraic Stack

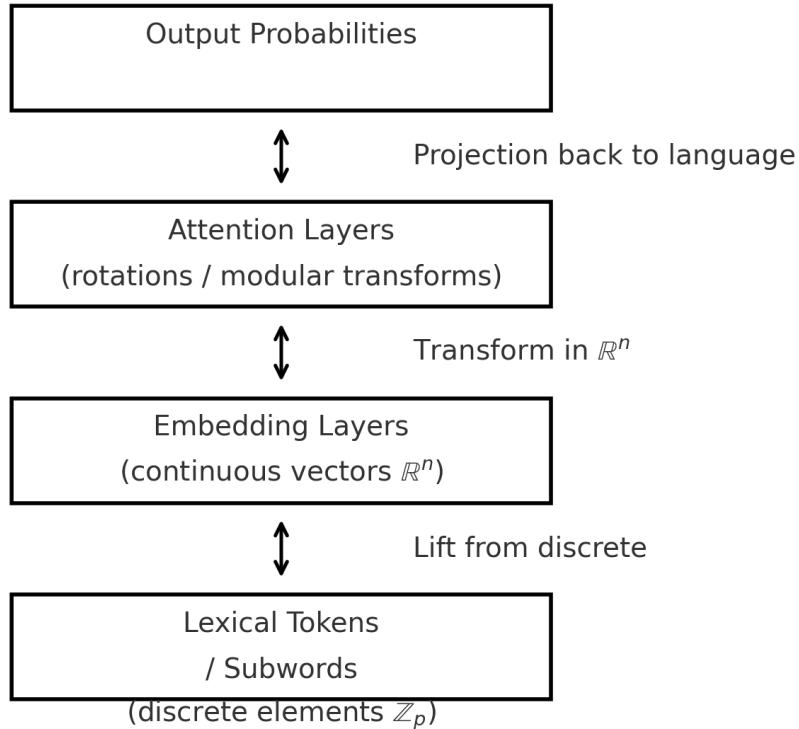


Figure 3. The Algebraic Stack—Conceptual schematic of the representational hierarchy in a language model. Lexical tokens and subwords originate as discrete elements in the finite field \mathbb{Z}_p and are lifted into continuous vector space \mathbb{R}^n through embedding layers. Attention layers perform rotations and modular transforms within \mathbb{R}^n to capture compositional structure, and the output probability layer projects the transformed representation back into linguistic space.

• • •

Closing

By viewing language as algebra, we shift from seeing LLMs as opaque statistical systems to seeing them as **computational engines with internal, learnable logic**. Recognizing these parallels between algebra and cognition allows us to design models that **count, classify, and compute** with greater transparency.

Future essays will extend this framework toward **empirical visualization of token-space periodicity**.

• • •

Figures

Figure 1—The Tokenization → Algebra Pipeline

Figure 2—Attention as Modular Arithmetic

Figure 3—The Algebraic Stack: a vertical flow from Discrete Tokens

$(\mathbb{Z}_p) \rightarrow \text{Embeddings}(\mathbb{R}^n) \rightarrow \text{Attention ("rotations")} \rightarrow \text{Output Probabilities.}$

• • •

References

- Cohen & Welling (2016). *Group Equivariant Convolutional Networks*.
- Tancik et al. (2020). *Fourier Features Let Networks Learn High-Frequency Functions*.
- Bronstein et al. (2021). *Geometric Deep Learning: Grids, Groups, Graphs, Geodesics*.
- Token Topology: A Geometric View of Language Representations (2024).

Related Articles—The Trilogy

- 1 [What Finite Fields Can Teach Us About Building Smarter Language Models](#)
- 2 [How Modular Arithmetic Powers Everything from AI to Cryptography](#)
- 3 [Tokenization + Algebra: How Language Models Learn to Count, Classify, and Compute](#)

• • •

Tags

#AI #LanguageModels #Algebra #MachineLearning #Interpretability
#Tokenization #ModularArithmetic #Cognition